



Data X

About Me:

Introduction To Convolution Neural Networks

Ikhlaz Sidhu, Chief Scientist, SCET/IEOR, UC Berkeley
Sana Iqbal, GSI, Data-X, UC Berkeley

Review: A simple Feed Forward Neural Net

1. Accepts the Input.

2. Creates different combinations of **weighted input**

3. Induces **Non-linearity** in the input using Activations.

4. Repeats 1,2,3 in **hidden layers**.

5. Gives the Output.


Data^x



Review: Convolution functions in Image Processing

- **Convolution functions** applied between an image matrix and a filter matrix gives us information about the spatial distribution of the image features.
- Using **different types of filters** affects the information extracted from an image- blur, edges, texture.

Convolutional of Two Signals



$$f[x,y] * g[x,y] = \sum_{n_1=-\infty}^{\infty} \sum_{n_2=-\infty}^{\infty} f[n_1,n_2] \cdot g[x-n_1,y-n_2]$$

elementwise multiplication and sum


Data^x

Before We Go On

X is an image



h is filter (or kernel)

$$\begin{bmatrix} -1 & 0 & 0 \\ 0 & 0 & 0 \\ 0 & 0 & 1 \end{bmatrix}$$


What is $X * h$?


Data^X

Before We Go On

X is an image



h is filter (or kernel)

$$\begin{bmatrix} -1 & 0 & 0 \\ 0 & 0 & 0 \\ 0 & 0 & 1 \end{bmatrix}$$


What is $X * h$?



Answer:








1. It is another image
2. Depending on the kernel, it highlights specific features

Data^X

Many Operations with Convolution



These filter are more typical for image processing, but not feature extraction.

In CNNs, we are looking to extract features

Operation	Filter	Convolved Image
Identity	$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$	
Edge detection	$\begin{bmatrix} 1 & 0 & -1 \\ 0 & 0 & 0 \\ -1 & 0 & 1 \end{bmatrix}$	
	$\begin{bmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{bmatrix}$	
	$\begin{bmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{bmatrix}$	
Sharpen	$\begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix}$	
Box blur (normalized)	$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$	
Gaussian blur (approximation)	$\frac{1}{16} \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$	

Data^x

Feature Extraction



Visualization of the receptive field

0	0	0	0	0	0	30
0	0	0	0	50	50	50
0	0	0	20	50	0	0
0	0	0	50	50	0	0
0	0	0	50	50	0	0
0	0	0	50	50	0	0
0	0	0	50	50	0	0

Pixel representation of the receptive field

*

0	0	0	0	0	30	0
0	0	0	0	30	0	0
0	0	0	30	0	0	0
0	0	0	30	0	0	0
0	0	0	30	0	0	0
0	0	0	30	0	0	0
0	0	0	0	0	0	0

Pixel representation of filter

Multiplication and Summation = $(50*30)+(50*30)+(50*30)+(20*30)+(50*30) = 6600$ (A large number!)

Value of convolution at this spot is large because the sum of products

Think of this like a “moving dot product” of the kernel across the original image

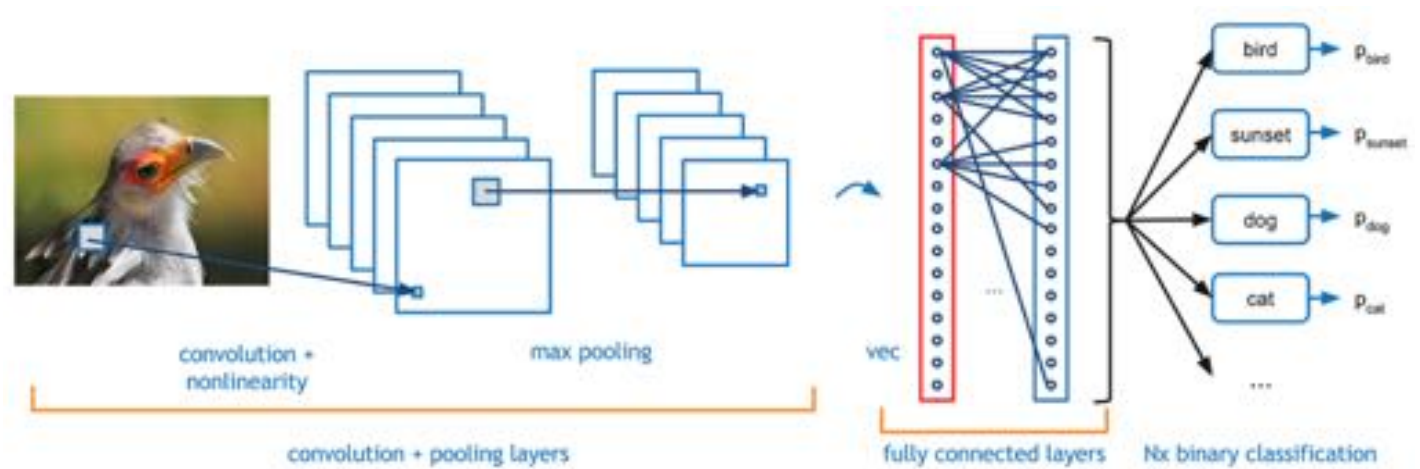
<https://adeshpande3.github.io/adeshpande3.github.io/A-Beginner%27s-Guide-To-Understanding-Convolutional-Neural-Networks/>

Data^x

Convolution Neural Networks

So there are four main **operations** in a basic CNN:

1. Convolution
2. Non Linearity (ReLU, tanh, sigmoid)
3. Pooling or Down-Sampling of features
4. Prediction

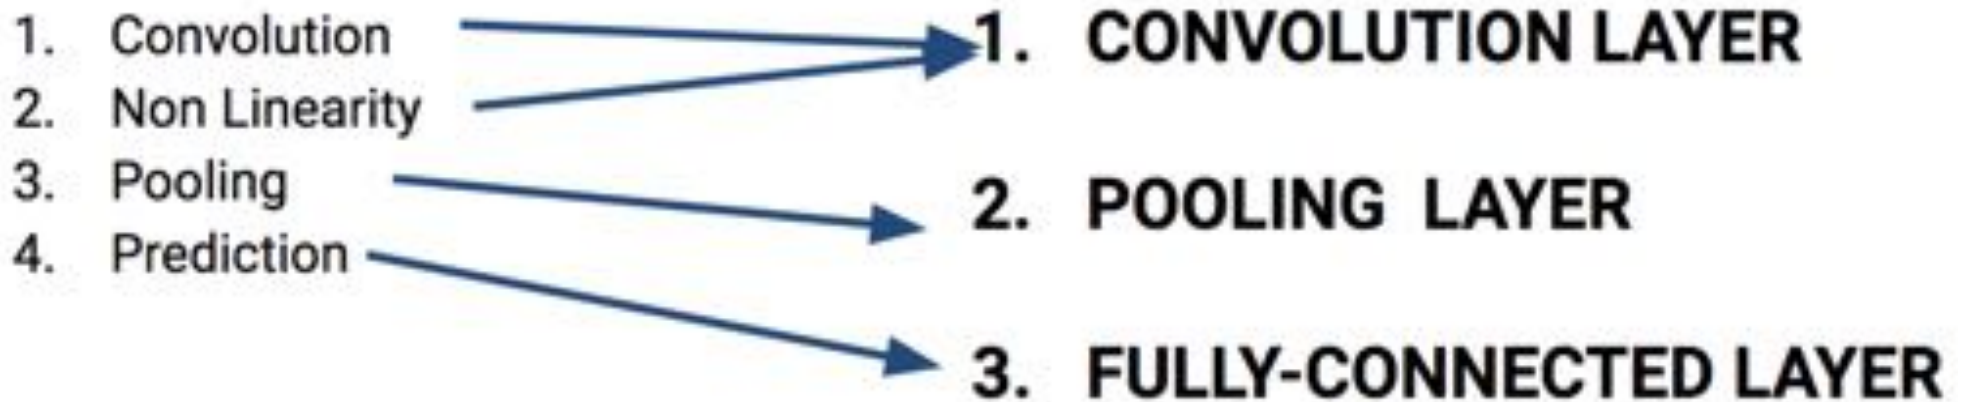


<https://adeshpande3.github.io/adeshpande3.github.io/A-Beginner%27s-Guide-To-Understanding-Convolutional-Neural-Networks/>

Data^x

Convolution Neural Networks

Three different **types of layers** in CNN:



Data^x

1. CONVOLUTION LAYER

- A. This layer has the filters that we want to apply on our image
- B. Each filter also called kernel has same depth as the channel of the image

For an image with **one channel** the below filter works as:

Sizes:

Image of $h \times w$

Kernel of $f \times g$

New image: $(h-f+1) \times (w-g+1)$

1	0	1
0	1	0
1	0	1

1 _{x1}	1 _{x0}	1 _{x1}	0	0
0 _{x0}	1 _{x1}	1 _{x0}	1	0
0 _{x1}	0 _{x0}	1 _{x1}	1	1
0	0	1	1	0
0	1	1	0	0

Image

4		


Convolved
Feature

Data^x


1. CONVOLUTION LAYER

Hyperparamters that can be tuned to change complexity and size of extracted features


Filter Size, **Stride** (step size of the filter), **Zero-padding** (to maintain dimensions)



- Stride 1
- No Zero padding
- Kernel size 3x3
- Input size 4x4
- Output size 2x2



- Stride 1
- Zero padding 1
- Kernel size 3x3
- Input size 5x5
- Output size 5x5



- Stride 2
- Zero padding 1
- Kernel size 3x3
- Input size 5x5
- Output size 3x3

$$n_{out} = \left\lfloor \frac{n_{in} + 2p - k}{s} \right\rfloor + 1$$

n_{in} : number of input features
 n_{out} : number of output features
 k : convolution kernel size
 p : convolution padding size
 s : convolution stride size


$n_{in} = n_{out}$ when

$$\text{Zero Padding} = \frac{(K - 1)}{2}$$

(K = kernel size)

1. CONVOLUTION LAYER


For an image with **3 channels** there would be **1 filter** for each channel



Data^x

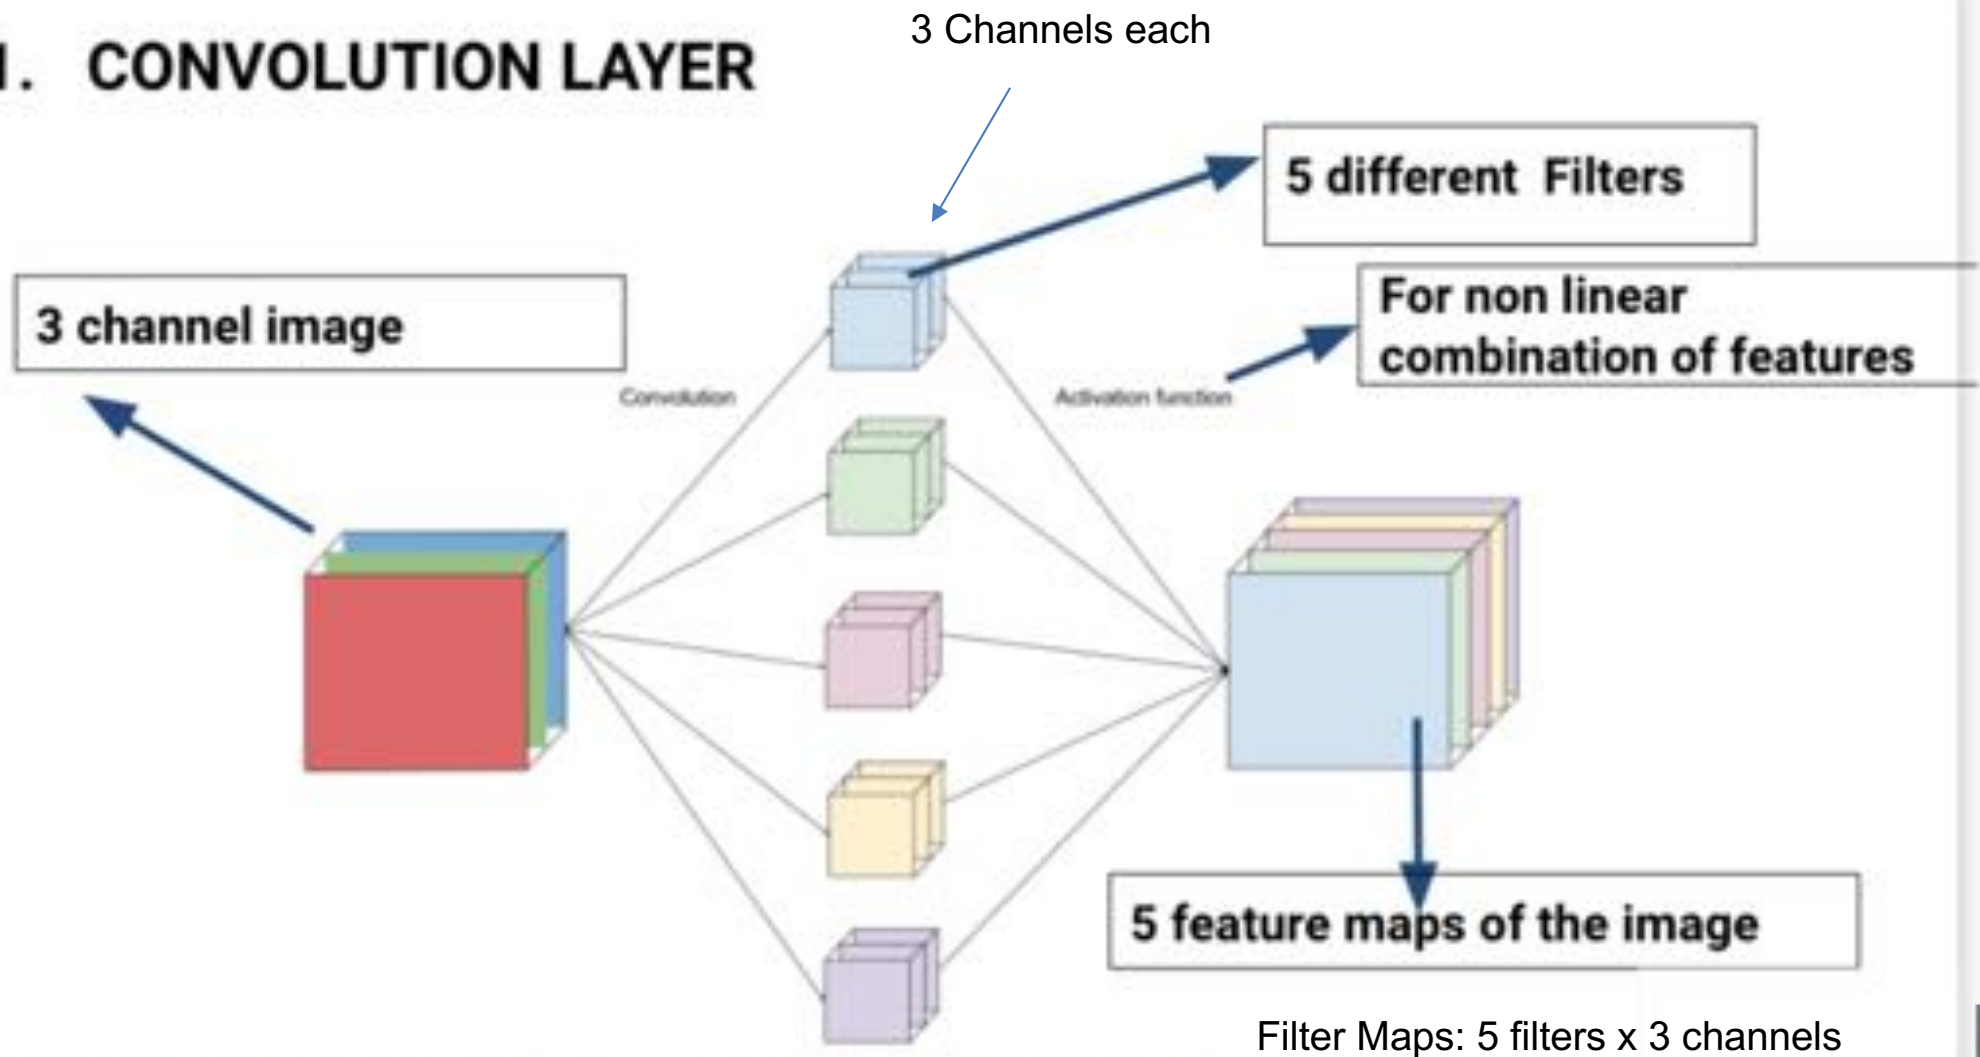
1. CONVOLUTION LAYER

For an image with 3 channels and multiple filters



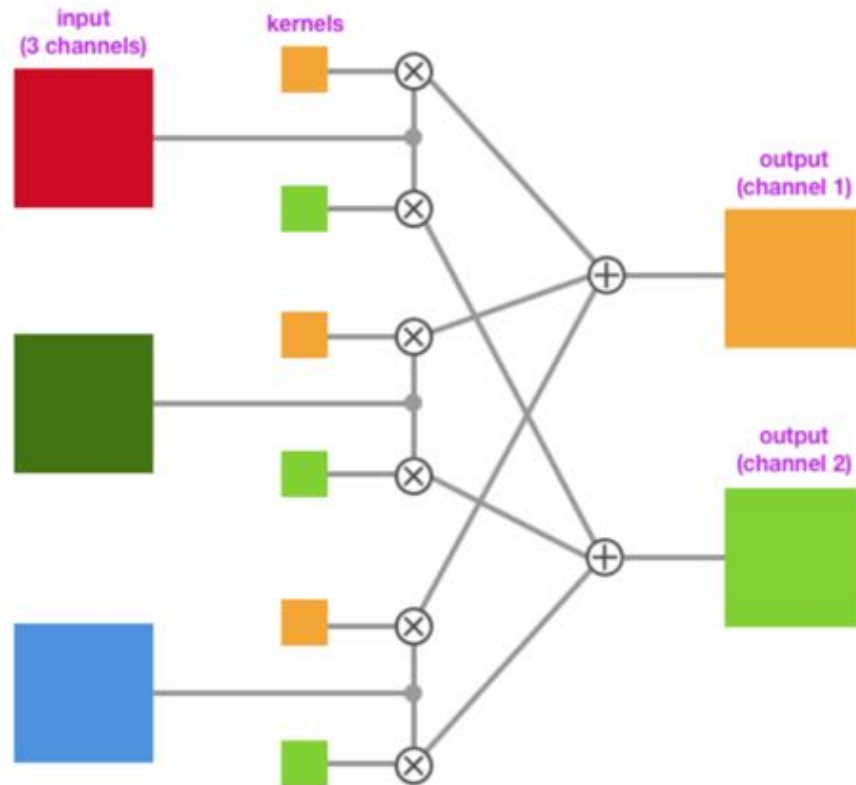
Data^x

1. CONVOLUTION LAYER





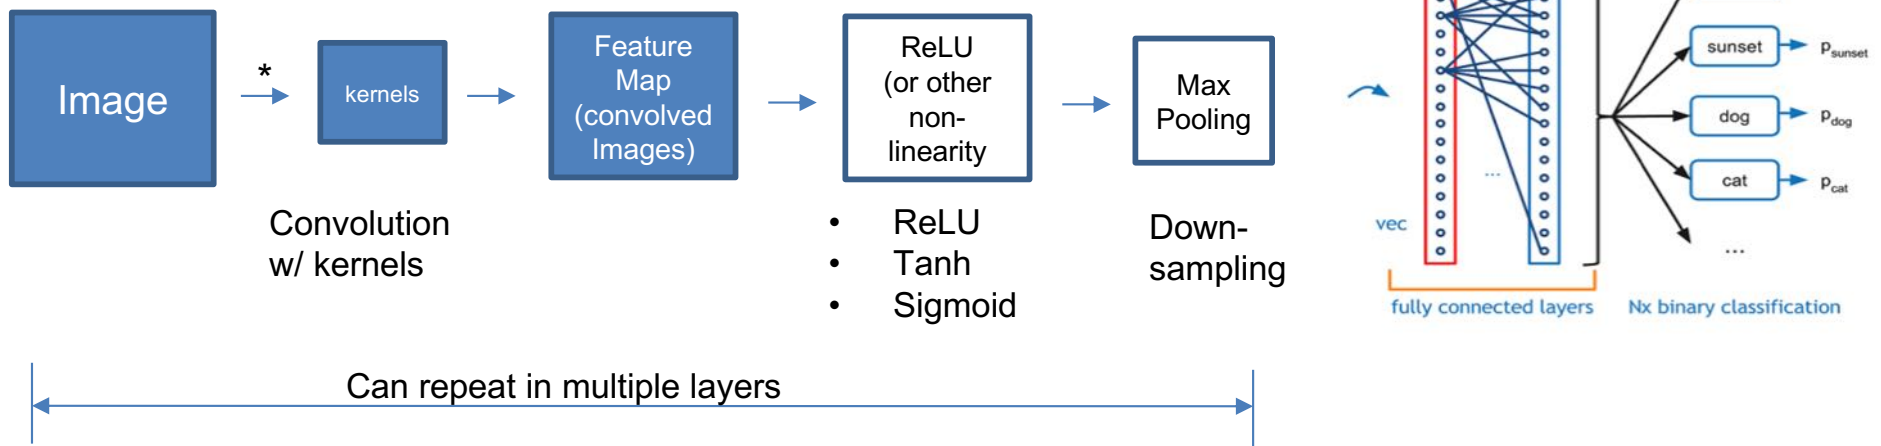
1. CONVOLUTION LAYER

A detail: Its is possible to sum the convolutions of 2 or more kernels



Data^x


- 1) Convolution and Non-linearity Blocks First
- 2) Then Fully Connected Layers leading to prediction



Data^x

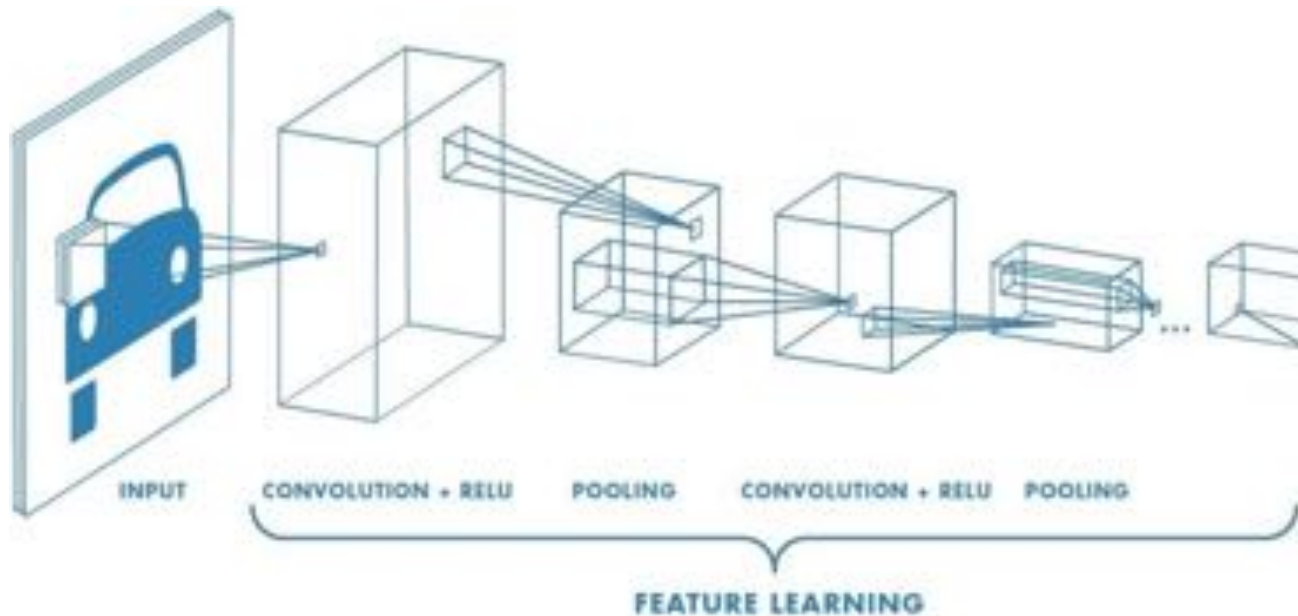
2. POOLING LAYER

1. Reduces the spatial dimensions (Width x Height) of the Input Volume
- 'down-sampling'
2. Usually max or average pooling is done.



Data^x

Putting the CNN Together



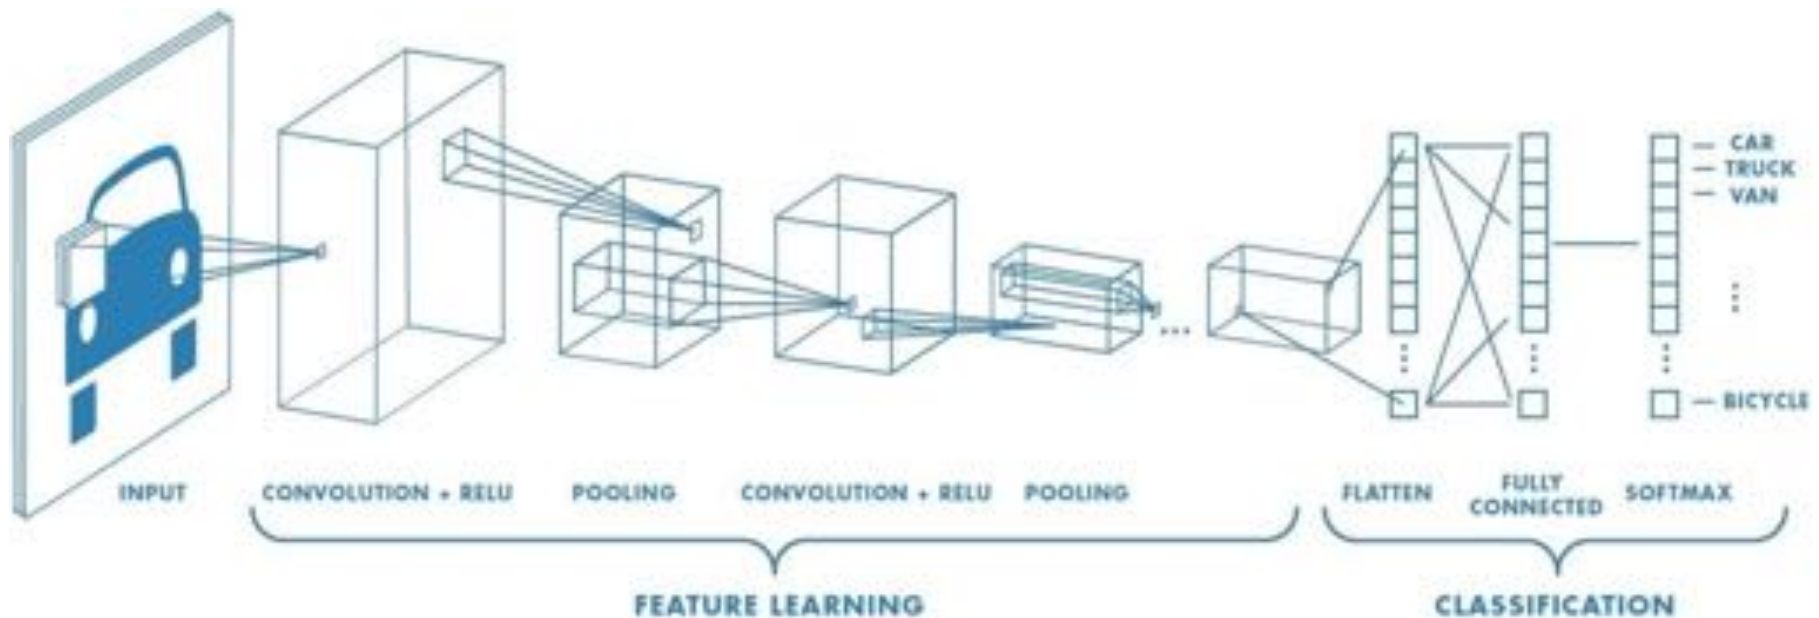
<https://www.mathworks.com/discovery/convolutional-neural-network.html>

Pooling:

- Reduces the spatial dimensions (Width x Height) of the Input Volume
- 'down-sampling'
- It does not affect the depth dimension of the Volume.
- Usually placed after the Convolutional layer.
- The decrease in size leads to less computational overhead for the upcoming layers of the network

Data^x

Putting the CNN Together with the Fully Connected Layer




<https://www.mathworks.com/discovery/convolutional-neural-network.html>

Data^x

What are the Parameters— What the Algorithm Chooses?

Putting the CNN together with the Fully Connected Layer



What are the Hyper-Parameters: What you choose in the CNN design?

<https://www.mathworks.com/discovery/convolutional-neural-network.html>

Data^x

Putting the CNN Together with the Fully Connected Layer What are the Parameters – What the Algorithm Chooses?

The filter or kernel itself!

What are the Hyper-Parameters: What you choose in the CNN design?

Number of kernels (features)

Sizes of the features

Number of Layers

Pooling Window Size, Pooling Stride

Fully Connected Layer: Number of Neurons, Number of Outputs

<https://www.mathworks.com/discovery/convolutional-neural-network.html>

Data^x

End of Section

Data^x